

Crystie Johnson

Lead SQL Developer / Data Architect

📍 Lawrenceville, GA 📞 470-377-1316 ✉️ BiDeveloperServices@gmail.com

🌐 [www.Linkedin.com/in/MrsCrystie](https://www.linkedin.com/in/MrsCrystie)

Github: <https://github.com/Crystie-Johnson>

US Citizen | Government for Public Trust

PROFESSIONAL SUMMARY

Accomplished Developer with 20+ years of experience designing and delivering enterprise-scale data platforms across Azure, AWS, and hybrid environments. Expert in cloud data architecture, large-scale migrations, ETL/ELT frameworks, and modern Lakehouse solutions.

Proven record of accomplishment transforming legacy systems into scalable, high-performance data ecosystems supporting analytics, BI, and mission-critical operations. Combines deep hands-on engineering expertise with strategic architecture leadership to deliver measurable business impact.

CORE COMPETENCIES

- Cloud Platforms – Azure, AWS, GCP (BigQuery)
- Master Data Management (MDM) – Customer, Product, Vendor Domains
- Data Architecture (EDW, Lakehouse, Data Warehousing, Data Lakes)
- ETL/ELT & Pipelines – Batch & Streaming (ADF, Airflow, DBT) Data Migration & Modernization
- Bronze / Silver / Gold Data Layering
- Data Modeling (Dimensional, Relational)
- SQL Performance Tuning & Optimization / Query Tuning
- Data Governance & Quality (MDM, CDC, Data Quality, Metadata)
- Data Security, Compliance (HIPAA, PII), and Resiliency
- DevOps & CI/CD (Azure DevOps, Git, Jenkins)
- Agile/Scrum Delivery & Cross-Functional Leadership

TECHNICAL SKILLS

Cloud & Data Platforms:

Azure (ADF, Synapse, Databricks), AWS (S3, Redshift, RDS, Glue), Snowflake

Data Engineering & Big Data:

Spark (PySpark, SQL), Kafka, Airflow, DBT, Delta Lake, Kafka, Spark Structured Streaming

Databases:

SQL Server, Oracle, PostgreSQL, MySQL, Teradata, SAP HANA

ETL / BI Tools:

ADF, SSIS, Informatica (PowerCenter), ODI, Wherescape, DataStage, Pentaho, Power BI

Programming:

SQL (T-SQL, PL/SQL, PL/pgSQL), Python, Shell/Bash

DevOps & Tools:

Azure DevOps, GitHub, Jenkins, ServiceNow

Key Achievements

- Migrated 30+ enterprise systems to Azure, improving scalability and performance
- Reduced deployment time by 80% through DevOps and CI/CD implementation
- Improved SQL performance by 30% on high-volume platforms
- Designed and optimized 25TB+ enterprise data warehouse architecture
- Built scalable ETL pipelines supporting batch, CDC, and real-time processing

PROFESSIONAL EXPERIENCE**Nuvia Business Solutions (Remote)****Lead Data Architect / Lead Data Engineer**

Aug 2024 – April. 2026

Served as Lead Data Architect and Technical Lead for enterprise-scale data modernization and cloud transformation initiatives supporting a 25TB+ analytics ecosystem. Led the design and implementation of scalable cloud-native data platforms, enterprise Master Data Management (MDM) solutions, metadata-driven ingestion frameworks, and modern ETL/ELT architecture across Azure and Snowflake environments.

Key Contributions & Achievements

- Led the modernization and migration of legacy enterprise platforms including Oracle, SQL Server, SAP, and on-premises data warehouses into scalable cloud-based analytics ecosystems, significantly improving scalability, reliability, and performance.
- Architected and implemented enterprise data ingestion frameworks using Azure Data Factory (ADF) and Snowflake to process structured, semi-structured, and unstructured data from Oracle, SQL Server, SAP, JSON, XML, Excel, and flat-file sources.
- Implemented enterprise-scale data deduplication and survivorship frameworks across large-volume datasets using deterministic and probabilistic matching techniques, including rule-based matching, fuzzy logic, hash comparisons, and confidence scoring models to improve data integrity, eliminate redundancy, and maintain high-quality trusted data within Azure Data Lake and Databricks Lakehouse environments.

- Designed and developed scalable ETL/ELT pipelines supporting batch, incremental, and metadata-driven ingestion patterns with advanced error handling, auditing, logging, and automated reconciliation processes.
- Developed and enforced enterprise data tagging and metadata management frameworks to classify, label, and annotate datasets based on sensitivity, source systems, business domains, lineage, and compliance requirements, improving data governance, discoverability, auditability, and regulatory adherence across cloud-based data platforms and Lakehouse environments.
- Directed end-to-end data migration lifecycle activities including data profiling, cleansing, transformation, validation, reconciliation, and production cutover support.
- Architected and implemented enterprise Master Data Management (MDM) solutions including data standardization, survivorship logic, de-duplication frameworks, golden record creation, and domain mastering across customer, vendor, and product domains.
- Defined and managed Master Critical Data Elements (MCDEs), metadata standards, lineage tracking, source-to-target mappings, and enterprise transformation logic across multiple integrated systems.
- Designed and implemented Bronze/Silver/Gold medallion architecture layers to support enterprise analytics, operational reporting, governance, and downstream business intelligence workloads.
- Developed reusable, parameterized ETL frameworks and modular pipeline components to improve maintainability, scalability, and deployment efficiency across enterprise projects.
- Implemented enterprise data quality, metadata management, lineage tracking, and governance frameworks to improve data integrity, compliance, and auditability.
- Designed dimensional data models including star and snowflake schemas to support enterprise reporting, analytics, and KPI-driven dashboards.
- Implemented CI/CD automation pipelines using Azure DevOps, Git, and automated deployment strategies, significantly reducing manual deployment effort and improving release efficiency.
- Performed cloud platform performance tuning and optimization initiatives to ensure high availability, scalability, resiliency, and efficient workload execution across enterprise data platforms.

Dassault Falcon Jet (Remote)

Lead Data Architect / Data Engineer

Jan 2023 – Aug 2024

Led enterprise data architecture and warehousing strategy supporting large-scale aerospace data systems.

Key Achievements:

- Migrated legacy SAP and Oracle data into SQL Server, reverse-engineering SAP, Wherescape Red, and Qlik systems to accurately map tables, transformations, business rules, and keys for the new SQL Server DataMart.
- Designed, developed, and maintained scalable enterprise data ingestion pipelines to onboard structured, semi-structured, and unstructured data from batch and real-time streaming sources, including REST APIs, relational databases, flat files, XML/JSON datasets, and message queue platforms into Azure and Databricks analytics environments using Azure Data Factory, Databricks, Apache Spark/PySpark, and Delta Lake technologies.
- Assisted with the operationalization, deployment, and ongoing support of cloud-based ETL and data platform services by standardizing and automating operational processes, developing reusable workflows, creating technical documentation and knowledge base articles, and providing guidance and support to operations teams transitioning to cloud technologies and modern data engineering practices.
- Integrated ServiceNow with ETL/BI systems to automate alerts and trigger workflows across tools such as Azure Data Factory and SSIS.
- Utilized Azure DevOps and Git for version control, implementing branching strategies and pull requests to streamline deployments across Dev, QA, Staging, and Production.
- Defined and executed comprehensive cutover and reconciliation strategies across legacy systems, extraction, conversion, load files, and S/4 environments—covering Master and Transactional Data, financial balancing, and inventory quantity/value validation.
- Designed Aerospace Inventory DataMart enabling advanced analytics and reporting
- Designed and implemented scalable ETL pipelines using Azure Data Factory, SSIS, and Wherescape RED, supporting CDC, batch, and incremental load patterns; enabled efficient handling of high-volume data and multiple ingestion strategies to ensure timely and reliable updates to downstream systems.
- Built scalable ETL pipelines using Azure Data Factory, SSIS, and Wherescape Red, supporting CDC, batch, and incremental loading patterns. Ensuring workflows were capable of handling large and growing data volumes and supported multiple data ingestion strategies, including full batch loads, incremental updates, and change data capture, to efficiently keep downstream systems up to date.
- Integrated dbt with Azure Data Factory and CI/CD pipelines (GitHub Actions/Azure DevOps) for automated deployments and robust transformation workflows.
- Built parameterized, reusable SSIS and ADF ETL frameworks with advanced error handling and logging.
- Developed advanced stored procedures, triggers, and user-defined functions to support critical business processes.
- Improved query performance by **30%** through optimization and indexing strategies
- Integrated dbt with CI/CD pipelines to automate transformation workflows

- Led cross-functional teams and mentored engineers on architecture and performance tuning
- Collaborated with stakeholders to translate business requirements into scalable data solutions

Fortna (MHS Global)

Senior Data Engineer / Data Architect

Oct 2021 – Jan 2023

Led enterprise data architecture and warehousing strategy supporting large-scale aerospace data systems.

Key Achievements:

- Led migration of 30+ legacy systems to Azure, improving system performance and scalability
- Designed Databricks Lakehouse architecture using Delta Lake and PySpark
- Built scalable ELT pipelines using Databricks, DBT and Azure Data Factory
- Reduced deployment time by 80% through DevOps automation
- Created and maintained custom database objects, stored procedures, and ETL processes using SQL Server, T-SQL, XML, and Infor BODS.
- Implemented data quality validation using dbt testing frameworks
- Developed real-time and batch pipelines supporting large-scale analytics workloads
- Designed enterprise data warehouse and DataMart architectures with optimized schemas and hierarchies
- Built robust ETL/ELT pipelines (ADF, SSIS, dbt) supporting batch and incremental processing
- Implemented Change Data Capture (CDC) and delta processing for optimized data synchronization and historical tracking.
- Developed authoritative data models and schemas for operational and analytical
- Led data governance initiatives, including lineage tracking and validation frameworks
- Improved query performance by 30% through SQL tuning and indexing strategies
- Collaborated with stakeholders to translate business requirements into scalable data solutions

CDC – National Syndromic Surveillance Program (NSSP)

Senior Data Engineer / Data Architect

Mar 2020 – Oct 2021

Led the design, implementation, and support of mission-critical syndromic surveillance solutions for the National Syndromic Surveillance Program (NSSP) in a cloud-based Agile (Kanban/Scrum) environment. Developed scalable ETL frameworks, data warehouse

solutions, and real-time analytics to enhance public health reporting, system reliability, and nationwide data visibility.

Key Achievements:

- Designed and implemented ETL pipelines supporting national public health systems
- Built real-time data processing solutions using AWS (S3, Redshift, RDS, EC2)
- Transformed and loaded healthcare data in compliance with HL7 messaging standards and HIPAA regulations, ensuring secure handling of PHI, PII across ETL workflows
- Designed, implemented, and optimized complex SQL Server and Oracle database objects including schemas, indexes (Columnstore, clustered), stored procedures, triggers, views, and functions
- Implemented CDC and SCD models for historical data tracking
- Extracted and processed heterogeneous data formats such as EBCDIC, ASCII, XML, HL7, PII, JSON, and unstructured datasets using SSIS, Informatica, Pentaho, and DataStage. Integrated disparate data formats including large-scale CSV files, applying robust data validation, error handling, and cleansing routines to ensure high data quality and reliability
- Utilized batch files (bat), BCP, and shell scripting to automate ingestion workflows across platforms.
- Supported real-time reporting and analytics for nationwide health monitoring

JP Business Solutions, Inc.

Senior Data Architect / Data Engineering Consultant

Nov 1995 – Mar 2020

Provided enterprise data architecture, engineering, and modernization consulting services for Fortune 500 organizations, healthcare enterprises, financial institutions, telecommunications companies, logistics providers, and government agencies, including Comcast, Coca-Cola Enterprises, Northrop Grumman, UPS, BBVA Bank, Bank of America, Kaiser Permanente, Children's Hospital, BellSouth, and other large-scale enterprise clients.

Led enterprise-wide data transformation initiatives focused on cloud modernization, enterprise data warehousing, ETL/ELT architecture, advanced analytics, and scalable data integration frameworks across Azure, AWS, and hybrid enterprise environments.

Key Contributions & Achievements

- Delivered enterprise-scale data architecture and engineering solutions supporting mission-critical analytics, operational reporting, regulatory compliance, and executive decision-making for Fortune 500 and government clients.

- Designed and implemented large-scale enterprise data warehouse architectures, operational data stores (ODS), and scalable ETL/ELT frameworks supporting high-volume transactional and analytical workloads.
- Led cloud migration and modernization initiatives across Azure and AWS, successfully migrating enterprise applications, databases, and legacy data warehouse platforms into scalable cloud-native analytics ecosystems.
- Directed end-to-end ETL development using Informatica PowerCenter, developing scalable mappings, workflows, sessions, and reusable transformation components for complex enterprise data integration environments.
- Architected metadata-driven ETL frameworks to standardize ingestion, transformation, validation, and loading processes, significantly improving development efficiency, maintainability, and operational consistency.
- Managed workflow orchestration, scheduling, dependency management, restartability, and failure recovery processes using enterprise workload automation and Informatica PowerCenter Workflow Manager (PWC).
- Developed and optimized large-scale data processing pipelines utilizing Apache Spark, PySpark, Delta Lake, Kafka, Databricks Auto Loader, and MLflow to support both batch and near-time analytics use cases.
- Designed and implemented scalable ingestion frameworks integrating Azure Data Lake Storage (ADLS Gen2), Azure SQL, Blob Storage, Databricks, REST APIs, flat files, and enterprise relational databases.
- Built scalable dimensional and relational data models including star schema, snowflake schema, and third normal form (3NF) models supporting enterprise reporting, BI dashboards, KPI metrics, and advanced analytics.
- Performed advanced SQL optimization, indexing strategies, query tuning, and database performance enhancements to improve scalability and reduce processing times across large-scale enterprise systems.
- Implemented enterprise data governance, lineage tracking, metadata management, and documentation standards using Atlassian Confluence and related collaboration platforms.
- Documented enterprise ETL frameworks, Bronze/Silver/Gold medallion architectures, system integrations, source-to-target mappings, and end-to-end data flows to support operational transparency and governance initiatives.

EDUCATION

Bachelor of Science – Computer Science 1991

University of Georgia